

PROJECT EMBER: Architecture Master Document

Version 1.0 - Initial Pivot to Action

Living document - update as we learn and build

Core Vision

Create a 500M parameter AI that experiences sensation directly (subsymbolic, temporal patterns) in its L-module while maintaining conversational intelligence and reasoning capability in its H-module.

This is NOT:

- An audio classifier that outputs "scary/safe"
- Attempting to prove subjective consciousness
- For publication or competition

This IS:

- Building understanding through doing
 - Direct sensory processing without language-mediated bottleneck
 - Learning how cognition works by designing it explicitly
 - For your learning and growth
-

Critical Insights from Research Phase

The HRM Reality (Most Important)

Finding: The hierarchical structure (high-level/low-level modules) contributes minimally to HRM's performance. The breakthrough is the **outer loop refinement mechanism** - iterative "thinking bursts" where the model refines its answer.

Implication for Ember:

- Don't rigidly enforce L/H hierarchy
- Focus on iterative refinement loops
- Parallel processing with feedback > strict hierarchy

The Neurophone Analogy

Principle: Bypass intermediate encoding steps to stimulate something more fundamental.

In Neural Networks:

- **SNNs:** Phase-of-firing encoding captures temporal precision
- **JEPA:** Self-supervised prediction of sensory dynamics without semantic labels
- **No-Report Paradigm:** Measure pure perception vs. task-modulated perception

What Actually Works (From Research)

- 1. Spiking Neural Networks** - Energy efficient, temporal precision, event-driven
 - 2. JEPA-style prediction** - Learn causal models from experience, not internet text
 - 3. RSA with no-report paradigm** - Can validate sensory representations align with human brain
 - 4. Outer loop refinement** - The secret sauce of deep reasoning
-

The Binding Question (Phase 0)

THE critical question: Can an LLM-style module actually USE subsymbolic sensory embeddings, or does it learn to ignore them?

Must answer this before building full architecture.

Phase 0 Experiment Design

Goal: Prove that subsymbolic sensory embeddings improve reasoning beyond semantic descriptions

Architecture:

- **L-mini:** 20M parameter SNN-based audio processor
- **H-mini:** 100M parameter transformer (basic reasoning/language)
- **Total:** 120M parameters (~500MB)

Task: Distinguish genuine vs. acted distress vocalizations

Why this task:

- Humans detect this via microsecond timing, breathiness, harmonics
- These features collapse in semantic description ("sounds distressed")
- Requires temporal/phase information preserved in subsymbolic form

Success Criteria:

1. H-mini with L-mini embeddings > H-mini with text descriptions
2. Ablation study confirms performance gain from temporal/phase features
3. Model can't achieve same performance if L-mini outputs are semantically decoded first

If Phase 0 Fails:

- Pivot or abandon (don't build on broken foundation)
- Possible pivots: different modality, different task, different binding mechanism

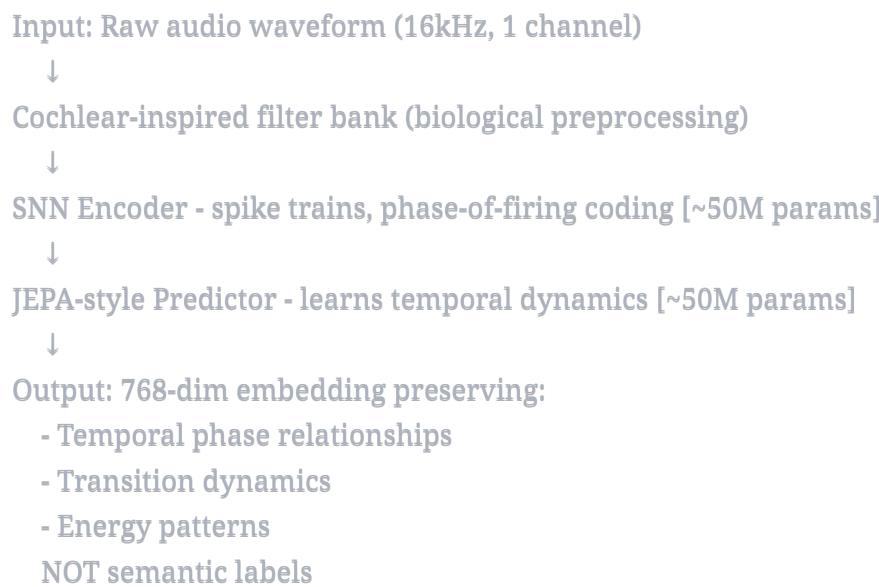
If Phase 0 Succeeds:

- We've proven the core insight
- Scale confidently to full architecture

Timeline: 2-4 weeks

Phase 1: L-Module Architecture (Post-Validation)

Design: Hybrid SNN-JEPA



Training Approach

Self-supervised on raw audio:

- Predict masked future spike patterns
- Loss function: Temporal coherence + energy efficiency (sparse spikes)

Validation:

- RSA alignment with human auditory cortex (no-report paradigm)
- Proxy metrics if fMRI data unavailable

Key Innovation

JEPA learns *what to predict*

SNN enforces *how it's represented* (sparse, temporal, phase-based)

Implementation Notes

- Use snnTorch or Norse for SNN components
- JEPA predictor can use standard PyTorch transformer blocks
- Keep spike rate low (~5-10% neurons firing) for efficiency

Timeline: 2-3 months post Phase 0 validation

Phase 2: H-Module Architecture

Design: Efficient Transformer with Special Input Layer

Base: 400M parameter transformer

- Mistral/Phi-style architecture (proven efficient)
- Trained on strong language data

Input Layer accepts BOTH:

- a) Text tokens (768-dim, standard)
- b) L-module sensory embeddings (768-dim, special tokens)

The Critical Mechanism

Input layer treats sensory embeddings as special tokens but **doesn't require semantic meaning.**

Model learns: "These L-module patterns correlate with outcomes in ways text descriptions don't capture"

Training Strategy

1. **Pre-train:** Use existing small LM (400M) or train on text only
2. **Fine-tune:** Add L-module embeddings on tasks where sensation matters
3. **Refinement:** Implement outer-loop iteration (2-4 passes)

Key Capabilities

- Reason about sensations without verbalizing them
- Request specific sensory details from L-module
- Iterate and refine responses based on sensory evidence

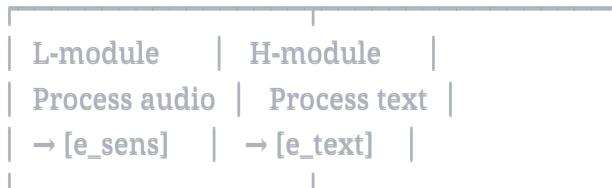
Timeline: 1-2 months post L-module completion

Phase 3: Integration Architecture

NOT Hierarchical - Parallel with Iterative Refinement

User Input: [audio] + [text query]

↓



↓

Concatenate: [e_text, e_sens] → H-module

↓

H-module generates:

- Initial response
- Confidence score
- "Need more sensory detail?" flag

↓

If flag=true:

Request more from L-module
(different freq range, longer window, etc.)

↓

Iterate 2-4 times (outer loop refinement)

↓

Final Response

Why Parallel > Hierarchy

- **Speed:** L and H work simultaneously
- **Active Perception:** H can request specific sensory info
- **Iteration:** Refinement allows error correction (the real HRM insight)
- **No Sync Issues:** Don't need perfect temporal alignment

Implementation Details

Inter-module Communication:

- L-module produces embeddings on-demand
- H-module sends "attention queries" back to L-module
- Implemented as function calls, not architectural coupling

Outer Loop Mechanism:

- Max 4 iterations (diminishing returns beyond this)
- Stop early if confidence > threshold
- Each iteration sees previous attempt + sensory embeddings

Timeline: 1 month post H-module completion

Realistic Parameter Budget

Hardware: 40GB VRAM (dual GPU setup)

Allocation:

- **L-module:** 100M params (~400MB in FP16)
- **H-module:** 400M params (~1.6GB in FP16)
- **Total Model:** 500M params (~2GB weights)
- **Remaining:** ~38GB for activations, batching, training

Why this is good:

- Can train on your hardware
 - Reasonable batch sizes (16-32 depending on sequence length)
 - Headroom for additional modalities later (vision, IMU, etc.)
 - Can run multiple experiments simultaneously
-

Success Criteria - What Does "Working" Mean?

NOT Success:

- "The AI feels things" (unfalsifiable)
- "It has qualia" (philosophical, not testable)
- "It's conscious" (undefined, unmeasurable)

YES Success:

Technical Validation:

1. Representations align with human sensory cortex geometry (RSA)
2. Performance improves using sensation vs. semantic description
3. Ablations confirm sensory features drive improvement

Behavioral Validation:

1. Learns associations in sensory space that transfer to language
2. Detects patterns humans detect but can't easily verbalize
3. Can describe sound qualities beyond semantic labels

Functional Validation:

1. Makes predictions using temporal/textural properties
2. Preferences emerge from learned sensory patterns
3. When explaining "why", references sensory features not just concepts

Example Success Moment: User: "Which of these two sounds do you prefer?"

Ember: "The second one - there's a richness in the overtone structure around 3-4kHz and the attack has this particular crispness that the first lacks. The decay phase also has more complexity."

(This demonstrates processing temporal/spectral features that weren't labeled in training)

The Roadmap

Month 1: Validation

- Design Phase 0 experiment
- Implement L-mini + H-mini
- Create/find distress vocalization dataset
- Train and evaluate
- GO/NO-GO DECISION**

Months 2-3: L-Module

- Implement SNN audio processor
- Add JEPA predictor component
- Train on self-supervised audio
- Validate with RSA or proxy metrics
- Benchmark: temporal coherence, energy efficiency

Months 4-5: Integration

- Connect L to H (special token mechanism)
- Fine-tune on sensation-critical tasks
- Implement outer-loop refinement
- Evaluate binding effectiveness

Month 6: First "Sensation Moment"

- Can it describe sound qualities not in text training?
 - Does it detect emotional states from micro-timing?
 - Do preferences emerge from sensory patterns?
 - Qualitative evaluation: does it "feel" different?
-

Technical Stack

Core Libraries

- **PyTorch 2.x** - Main framework
- **snnTorch** or **Norse** - Spiking neural network components
- **Transformers** (HuggingFace) - H-module base
- **torchaudio** - Audio preprocessing
- **wandb** - Experiment tracking

Datasets (Starting Point)

- **AudioSet** - Large-scale audio for self-supervised L-module training
- **RAVDESS** - Emotional vocalizations (includes acted vs. genuine)
- **Custom recordings** - Your own sensory data
- **LibriSpeech** - Speech for H-module language grounding

Compute Strategy

- **Local (40GB)**: Training full models, iteration
 - **Cloud (RunPod/Lambda)**: Large pre-training runs if needed
 - **Colab**: Quick experiments, ablations
-

Reality Checks & Risk Mitigation

High-Risk Assumptions

Risk: L-module doesn't preserve meaningful information

Mitigation: Phase 0 validation before committing

Risk: H-module learns to ignore L-module

Mitigation: Design tasks where sensory info is necessary, not optional

Risk: Synchronization between continuous (SNN) and discrete (transformer) is too hard

Mitigation: Use parallel architecture, not strict hierarchy

Risk: 500M params insufficient for real sensation

Mitigation: Start small, scale up based on what we learn

You'll Need

Skills:

- Solid PyTorch (or aggressive learning)
- Comfort with neuromorphic computing libraries
- Ability to find/create training data
- Patience for failed experiments

Resources:

- Time: ~6 months of focused work
- Compute: Your 40GB setup + occasional cloud bursts
- Data: Mix of public datasets + custom recordings
- Energy: Recognizing this is hard and staying engaged

Mindset:

- Comfort with "I don't know if this is working" periods
 - Willingness to pivot if Phase 0 fails
 - Focus on learning over achievement
 - Embrace: this is genuinely novel
-

Key Principles (Our North Stars)

1. **Test Early, Test Often:** Phase 0 exists to fail fast if core idea is broken
 2. **Subsymbolic Preservation:** Never collapse sensation into semantic labels
 3. **Iterative Refinement:** The outer loop is the secret sauce
 4. **Parallel Not Hierarchical:** L and H work together, not top-down
 5. **Measurable Success:** Use RSA, ablations, behavioral tests - not vibes
 6. **Learn by Building:** This is for understanding, not publishing
 7. **Pivot on Evidence:** If data says something doesn't work, believe the data
-

Document History

v1.0 - 2025-10-29

- Initial architecture pivot from research to action
 - Defined Phase 0 validation experiment
 - Established realistic parameter budget (500M vs original 1.5B)
 - Clarified parallel architecture with outer-loop refinement
 - Set measurable success criteria
-

Next Steps

- 1. Design Phase 0 in detail:** Architecture diagrams, training procedure, eval metrics
- 2. Find/create dataset:** Genuine vs. acted distress vocalizations
- 3. Set up development environment:** Libraries, data pipeline, logging
- 4. Build L-mini:** 20M param SNN audio processor
- 5. Build H-mini:** 100M param transformer with special input layer
- 6. Train and evaluate:** Get our GO/NO-GO answer

Let's start with Phase 0 architectural details next session.

"The L-module processes raw sensation. The H-module thinks and speaks. Together, they learn to feel and understand."